

CSCI 5673

Distributed Systems

Lecture Set Zero

Introduction

Lecture Notes by
Shivakant Mishra
Computer Science, CU Boulder
Last update: January 17, 2017

Introduction


- What is a distributed system?
 - A distributed system is a collection of independent computing components, interconnected via a network, that are capable of collaborating on a task
 - Independent computing components
 - Each computing component has its own set of resources such as CPU, memory, disk, etc and can execute on its own
 - Interconnected: computing components can exchange information using the underlying network

Distributed Systems: Key Features

- No centralized control
- An integration of system services, presenting a transparent view of a multiple computer system with distributed control and resources

Introduction

- Distributed Systems is NOT...
 - Advanced Networking
 - Hardware routing architectures
 - Optical networking
 - TCP enhancements
 - Advanced Operating Systems
 - Multi-core systems – scaling an OS to 128 processors
 - tightly coupled vs. loosely coupled systems
 - Microkernels, etc.



Introduction

- Distributed Systems is NOT...
 - Network Security
 - Distributed Denial of Service (DDOS) attacks
 - Stuxnet virus
 - Rootkits
 - Advanced cryptography, ETC...
 - Software Engineering
 - Thread-based implementations of servers

Scope

- Assumptions about what you know:
 - Operating Systems
 - Processes, threads, synchronization, mem mgmt/virtual memory, file systems,
 - Distributed file systems, VMs
 - Network Systems
 - Ethernet/WiFi/MAC protocols, VLANs
 - IP routing, RIP, OSPF, BGP, IP multicast, MPLS
 - TCP/UDP transport layers
 - SSL/SSH, public/symm key, certificates, VPNs
 - HTTP, DNS, SMTP application layer protocols
 - Socket API and programming

Why distributed systems?

- Price/performance ratio
 - Don't get twice computing for twice the price with multiprocessors
- Distributing applications may make sense
 - ATMs, graphics, remote monitoring
- Interactive communication & entertainment
 - work and play together:
email, gaming, telephony, instant messaging
- Remote content
 - web browsing, music & video downloads, IPTV, file servers

Why distributed systems?

- Mobility
- Increased reliability
- Increased scalability
- Incremental growth

Problems

Designing distributed software is difficult

- Operating systems handling distribution
- Programming languages?
- Efficiency?
- Reliability?
- Administration?

Network

- disconnect, loss of data, latency

Security

- want easy and convenient access

Limitations of Distributed Systems

- Absence of a common global clock
- Absence of shared memory

- Absence of a common global clock
 - Notion of a common global time does not exist
 - Each component has its own local clock
 - Local clocks may drift from one another
 - It is not possible to perfectly synchronize all local clocks because of non-deterministic communication and scheduling delays

- Absence of a shared memory
 - Notion of a common memory (RAM or disk) does not exist
 - Each component may have its own memory
 - No temporal relation among read/write operations in different memory
 - It is not possible to perfectly synchronize all read/write operations

Limitations of Distributed Systems

- Due to the absence of a common global time and a common shared memory
 - It is not always possible to reason about the temporal order of events in a distributed system
 - It is not always possible to collect an up-to-date state of the entire system