

# CSCI 5673

# Distributed Systems

## Lecture Set Four

### Asynchrony, Failure Models and Distributed Consensus

Lecture Notes by  
Shivakant Mishra  
Computer Science, CU Boulder  
Last update: February 14, 2017

# Synchrony and Failure Models

- Design of a distributed system is dependent on two important *system assumptions*
  - Synchrony Model
  - Failure model

# Synchrony Model

- Synchrony model refers to the assumptions about latencies in distributed systems
- Synchronous systems
  - There is a *known* bound on scheduling and communication delays
- Asynchronous systems
  - There is no *known* bound on communication or scheduling delays
- Timed asynchronous systems
  - There is no *known* bound on communication or scheduling delays, but a statistical model exists
    - e.g. 99% latencies are less than 2 msec

# Failure Models

- A failure model describes the behavior of a component when it fails
  - Fail stop failure
    - The component stops responding when it fails
    - No incorrect state transition
    - Detectable
  - Crash failure
    - The component stops responding when it fails
    - No incorrect state transition
    - Most common type of failure

- Performance failure

- (Correct) response is produced either too early or late.

- Omission failure

- The component omits to produce a response to one or more inputs.

- Byzantine failure

- Arbitrary failure
  - Incorrect state transition, coordinated failures, malicious intent, ...

Subset relationship among failure models

There are several other failure models proposed in literature

# Failure Models

Subset relationship among failure models

# Failure Models

There are several other failure models proposed in literature

# Distributed Systems

## Synchrony Model

- Generally, a distributed system designed for a synchronous environment is simpler
- Most non-critical distributed systems are designed for asynchronous environment

## Failure model

- Generally, a system designed to tolerate fail stop failures is simpler than a system designed to tolerate crash failure, ...
- Most non-critical systems are designed to tolerate crash failures



# Consensus Problem

- Fundamental problem in distributed systems
  - Mutual exclusion, replication consistency, message ordering, distributed locking, leader election are all instances of distributed agreement
- Problem
  - M processes:  $P: p_1, \dots, p_m$
  - Each process  $p_i$  stores a value  $v_i$
  - A subset  $F$  of  $P$  are faulty
  - Goal: Each process  $p_i$  calculates a consensus value  $a_i$ .

- Solution requirements

- Agreement: For every pair of non-faulty process  $p_i$  and  $p_j$ ,  $a_i = a_j$ . This is the consensus value.
- Validity: The consensus value is a function of the initial value  $v_i$  of all non-faulty processes
- Termination: All non-faulty processes eventually calculate a consensus value

Consensus algorithm depends on system model and failure model

# Distributed Agreement: Synchronous Systems

- Termination requirement: Agreement should be achieved with in  $r$  rounds of message exchanges for some fixed  $r$ .

# Byzantine Generals Problem

- Reference: Lamport, Shostak and Pease. The Byzantine Generals Problem. ACM TOPLAS, 4(3), 1982.
- Turkish sultan led an invasion into Byzantium empire
- Byzantium emperor has several armies, each led by a general
- Byzantine generals can survive (win or retreat and escape safely) if their action (attack/retreat) is coordinated
- Generals do not talk with one another directly
  - They send messengers
- Goal: Devise an algorithm that enables Byzantine generals to win or escape safely.


# Byzantine Generals Problem

- Turkish sultan led an invasion into Byzantium empire
- Byzantium emperor has several armies, each led by a general
- Byzantine generals can win (or retreat and escape safely) if their action (attack/retreat) is coordinated
- Generals do not talk with one another directly
  - They send messengers
- Goal: Devise an algorithm that enables Byzantine generals to win or escape safely.
- **Problem**:
  - *Some generals may be on Turkish payroll*
  - These treacherous generals will try to deceive (e.g. lie to) the loyal generals to prevent a coordinated action


# Byzantine Generals Problem

- Assumptions:
  - Synchronous, reliable channels
  - Messages cannot be altered
  - Receivers can authenticate the senders

- Result 1: Suppose that there are three generals and one of them is a traitor
  - Byzantine generals problem cannot be solved



L1: 1 attack; 1 retreat




L1: 1 attack; 1 retreat

**L1 cannot distinguish between the two scenarios**


- Result 2: Suppose that there are four generals and at most one of them is a traitor
  - Byzantine generals problem can be solved

### Algorithm

1. C sends order to everyone
2. Ls exchange the orders received with one another
3. Ls decide on an action based on majority



L1, L2, L3: 2 Attacks, 1 Retreat



L1, L2: 2 Attacks, 1 Retreat



- General result: Suppose that there are  $n$  generals and at most  $f$  of them are traitors
  - Byzantine generals problem can be solved, if  $n > 3f$

# Consensus

- Scenario 1:
  - Synchronous system
  - No process or message omission failures
- One-round algorithm:
  - Each process sends its value to all the processes.
  - If all values a process has (including its own) are 1 then decide 1. Otherwise decide 0.

*<Figure>*

# Consensus

- Scenario 2:
  - Synchronous system
  - No process failure
  - Message omission failure: Any number of messages may be lost.
- Solution: ???
- Theorem: There is no algorithm that solves the agreement problem for even two processes

Hint: What if all messages are lost?

Network partition

# Consensus

- Scenario 3:
  - Synchronous system
  - No messages are lost (No message omission failure).
  - Up to  $f$  processes may fail (fail stop failure)
- Algorithm: ???

# Consensus

- Scenario 3:
  - Synchronous system
  - Up to  $f$  processes may fail (fail stop failure)
  - No messages are lost (No message omission failure).
- Algorithm
  - Each process maintains a vector containing a value for each process  $\{0, 1, u\}$ .  $u$  = undefined.
  - One round:
    - Send your vector to all processes.
    - Update local vector according to the received vectors (in case local vector has a “ $u$ ”, and any of received vectors contain “0” or “1”).
  - After  $f+1$  rounds decide according to the local vector. If you have a majority of 1 in the vector then decide 1, otherwise decide 0.

- Scenario 3:

Homework: Construct an example to show that the processes may not reach a consensus with only  $f$  rounds, for  $n$  processes  $n > f$ , for any  $n$  and  $f$ .

# Consensus: Asynchronous Systems

- Scenario 1:
  - Asynchronous system
  - No process or message omission failures
- Algorithm: ???
- Algorithm: *similar to the synchronous case*
  - Each process sends its value to all the processes
  - After receiving a value from all processes, if all values a process has (including its own) are 1 then decide 1. Otherwise decide 0

# Distributed Agreement: Asynchronous Systems

- Asynchronous systems
  - There is no *known* bound on communication delays
- Scenario 1:
  - Asynchronous system
  - No process or message omission failures
- Algorithm: *similar to the synchronous case*
  - Each process sends its value to all the processes
  - After receiving a value from all processes, if all values a process has (including its own) are 1 then decide 1. Otherwise decide 0

Question: How is this algorithm different from the synchronous case?



# Distributed Agreement: Asynchronous Systems

- Scenario 2:
  - Asynchronous system
  - No process failure
  - Message omission failure: Any number of messages may be lost
- Algorithm: ???
- Algorithm:
  - There is no algorithm that solves the agreement problem for even two processes
    - Note that the scenario 2 of synchronous case is a special case of this scenario.

# Distributed Agreement: Asynchronous Systems

- Scenario 3:
  - Asynchronous system
  - Up to  $f$  processes may fail (fail stop failure)
  - No messages are lost (No message omission failure).
- Algorithm: ???

# FLP Result

- In an asynchronous distributed system, no algorithm can guarantee to reach a consensus between participating processes if one or more of them can fail by stopping
- Reference: M.J. Fischer, N. Lynch and M.S. Paterson. Impossibility of Distributed Consensus with One Faulty Process. JACM, 32(2), 1985.
- One of the most important result in distributed systems
- Definitively placed an upper bound on what it is possible to achieve with distributed processes in an asynchronous setting

## FLP result: Consequences

- It is impossible to distinguish between a process failure and a communication failure in an asynchronous distributed system
- It is impossible to reliably detect the failure of a process in an asynchronous distributed system

Question: Scenario 3 is a very common scenario in practice. So, what can we do in light of the FLP result?

# FLP Result: What does ``Impossibility'' mean?

- In formal proofs, an algorithm is totally correct, if
  - It computes the right thing
  - And it ``always'' terminates
- FLP proves that any algorithm that solves consensus in an asynchronous distributed system in the presence of process failures has runs that never terminate
  - These runs are extremely unlikely
- Consensus is impossible thus means consensus is not always possible

# Consensus: In Light of the FLP Result

- Probabilistic algorithms for consensus
  - Allow algorithms to not always guarantee consensus
- Allow algorithms to not always terminate
- Failure detectors
- Group membership protocols

# Safety and Liveness Properties

- Safety property: Nothing bad will ever happen
  - e.g. No two process will ever acquire a write lock for the same item at the same time
  - e.g. No two messages will ever be delivered in different order in different processes using a total order protocol
  - e.g. at no point, two or more new tokens will be generated
- Liveness property: Something good will eventually happen
  - e.g. a process requesting to acquire a lock will eventually acquire it
  - e.g. a message multicast using a total order atomic multicast protocol will eventually be delivered by all processes
  - e.g. a new token will eventually be generated after a token loss

- Distributed services built for an asynchronous environment where components may suffer crash failures can satisfy safety property but not liveness property
  - i.e. there will be some scenarios under which these services will not be able to provide an output; these scenarios are generally extremely unlikely