

# CSCI 5673

# Distributed Systems

## Lecture Set Eight

### Raft: An Understandable Consensus Algorithm

Lecture Notes by  
Shivakant Mishra  
Computer Science, CU Boulder  
Last update: February 28, 2017

D. Ongaro and J. Ousterhout. In Search of an Understandable Consensus Algorithm. USENIX ATC 2014.

# Motivation

- **Paxos**
  - Current standard for both teaching and implementing consensus algorithms
  - Very difficult to ***understand*** and very hard to ***implement***
- **Raft**
  - New protocol (2014)
  - Much easier to ***understand***
  - Several ***open-source implementations***

# Key features of Raft

- ***Strong leader:***
  - Leader does most of the work:
    - Issues ***all*** log updates
- ***Leader election:***
  - Uses ***randomized timers*** to elect leaders.
- ***Membership changes:***
  - New ***joint consensus*** approach where the majorities of two different configurations are required

# Raft consensus algorithm (I)

- Servers start by electing a ***leader***
  - Sole server habilitated to accept commands from clients
  - Will enter them in its log and forward them to other servers
  - Will tell them when it is safe to apply these log entries to their state machines

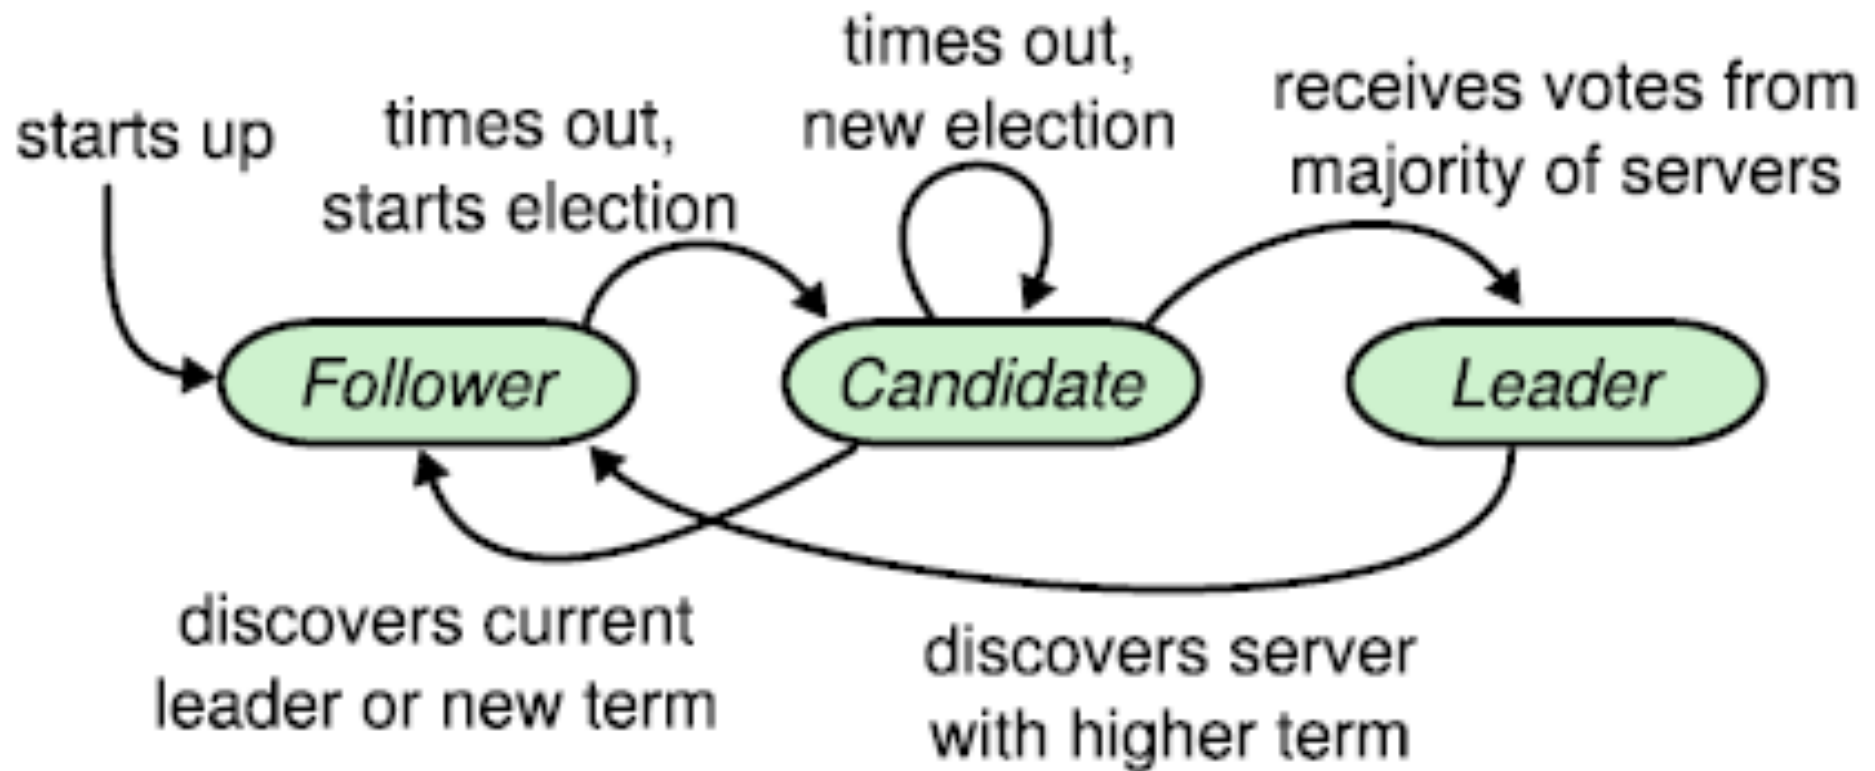
# Raft consensus algorithm (II)

- Decomposes the problem into three fairly independent subproblems
  - **Leader election:**  
How servers will pick a—**single**—leader
  - **Log replication:**  
How the leader will accept log entries from clients, propagate them to the other servers and ensure their logs remain in a consistent state
  - **Safety**

# Raft basics: the servers

- A RAFT cluster consists of several servers
  - Typically five
- Each server can be in one of three states
  - *Leader*
  - *Follower*
  - *Candidate* (to be the new leader)
- Followers are passive:
  - Simply reply to requests coming from their leader

# Server states

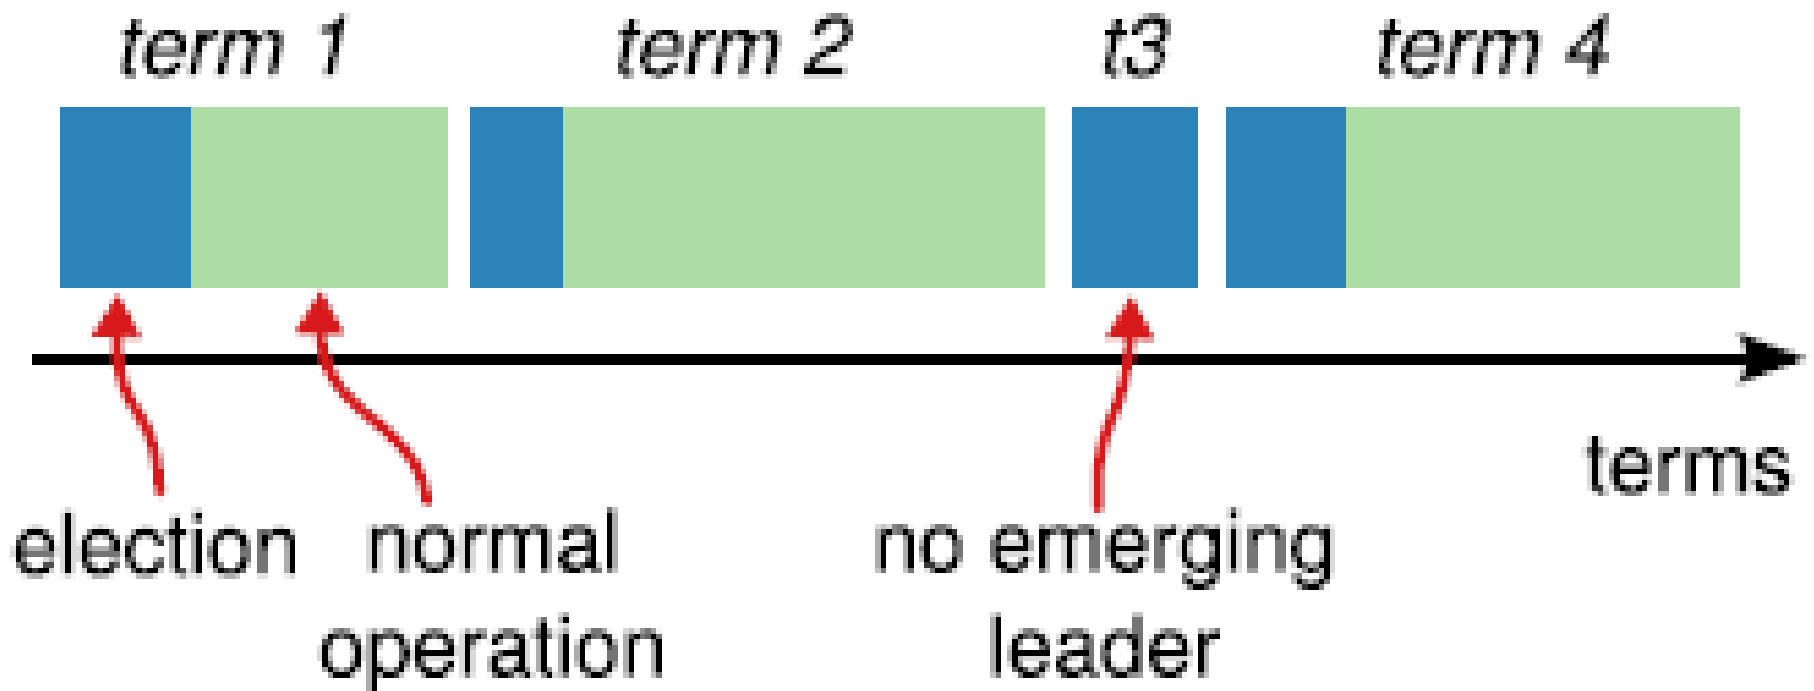




# Raft basics: terms (I)

- Epochs of arbitrary length
  - Start with the election of a leader
  - End when
    - No leader can be selected (split vote)
    - Leader becomes unavailable
- Different servers may observe transitions between terms at different times or even miss them

# Raft basics: terms (II)



# Raft basics: terms (III)

- Terms act as logical clocks
  - Allow servers to detect and discard obsolete information (messages from stale leaders, ...)
- Each server maintains a current term number
  - Includes it in all its communications
- A server receiving a message with a high number updates its own number
- A leader or a candidate receiving a message with a high number becomes a follower

# Raft basics: RPC

- Servers communicate through idempotent RPCs
  - **RequestVote**
    - Initiated by candidates during elections
  - **AppendEntry**
    - Initiated by leaders to
      - Replicate log entries
      - Provide a form of heartbeat
        - » Empty AppendEntry( ) calls

# Leader elections

- Servers start being *followers*
- Remain followers as long as they receive valid RPCs from a leader or candidate
- When a follower receives no communication over a period of time (the *election timeout*), it starts an election to pick a *new leader*

# Starting an election

- When a follower starts an election, it
  - Increments its current term
  - Transitions to candidate state
  - Votes for itself
  - Issues ***RequestVote*** RPCs in parallel to all the other servers in the cluster.

# Acting as a candidate

- A candidate remains in that state until
  - It wins the election
  - Another server becomes the new leader
  - A period of time goes by with no winner

# Winning an election

- Must receive votes from a majority of the servers in the cluster for the same term
  - Each server will vote for at most one candidate in a given term
    - The first one that contacted it
- Majority rule ensures that at most one candidate can win the election in a given term
- Winner becomes **leader** and sends heartbeat messages to all of the other servers
  - To assert its new role



# Hearing from other servers

- Candidates may receive an *AppendEntries* RPC from another server claiming to be leader
- If the leader's term is at greater than or equal to the candidate's current term, the candidate recognizes that leader and returns to follower state
- Otherwise the candidate ignores the RPC and remains a candidate

# Split elections

- No candidate obtains a majority of the votes in the servers in the cluster
- Each candidate will time out and start a new election
  - After incrementing its term number

# Avoiding split elections


- Raft uses randomized election timeouts
  - Chosen randomly from a fixed interval
- Increases the chances that a single follower will detect the loss of the leader before the others

# Reacp: Key features of Raft

- ***Strong leader:***
  - Leader does most of the work:
    - Issues ***all*** log updates
- ***Leader election:***
  - Uses ***randomized timers*** to elect leaders.
- ***Membership changes:***
  - New ***joint consensus*** approach where the majorities of two different configurations are required

# Example

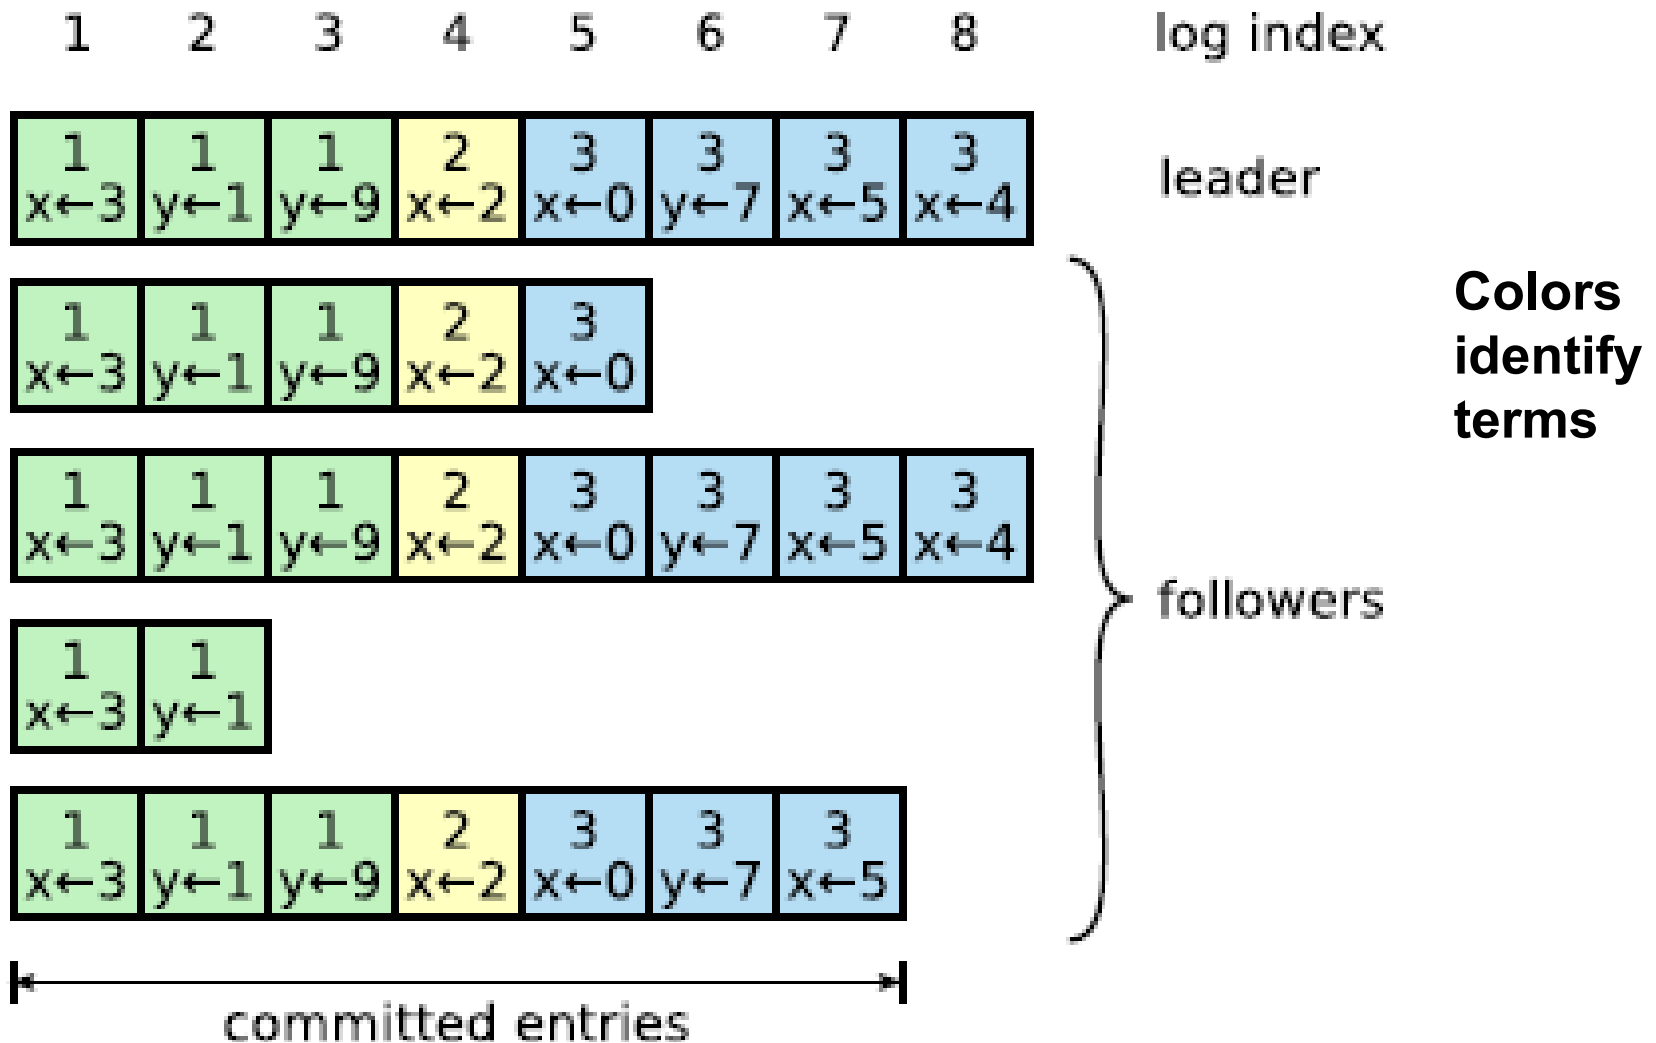
Follower with the *shortest timeout* becomes the *new leader*



# Log replication

- Leader
  - Accept client commands
  - Append them to their log (new entry)
  - Issue **AppendEntry** RPCs in parallel to all followers
  - Apply the entry to their state machine once it has been safely replicated
    - Entry is then ***committed***

# Log organization



# Handling slow followers ,...

- Leader *reissues* the AppendEntry RPC
  - They are idempotent



# Committed entries

- Guaranteed to be both
  - Durable
  - Eventually executed by all available state machines
- Committing an entry also commits all previous entries
  - All AppendEntry RPCs—including heartbeats—include the index of its most recently committed entry

# Why?

- Raft commits entries in ***strictly sequential order***
  - A log entry is committed once the leader that created the entry has replicated it on a majority of the servers
  - Requires followers to accept log entry appends in the same sequential order
    - ***Cannot "skip" entries***

**Greatly simplifies the protocol**

# Raft log matching property

- If two entries in different logs have the same index and term
  - These entries store the same command
  - ***All previous entries*** in the two logs are ***identical***


1	1	1	2	3	3	3	3
x←3	y←1	y←9	x←2	x←0	y←7	x←5	x←4

1	1
x←3	y←1

# Handling leader crashes (I)

- Can leave the cluster in a inconsistent state if the old leader had not fully replicated a previous entry
  - Some followers may have in their logs entries that the new leader does not have
  - Other followers may miss entries that the new leader has

# Handling leader crashes (II)



# Handling Crashes

- Missing entries: a-b
- Extra uncommitted entries: c-d
- Both: e-f

# Handling leader crashes (IV)

- Raft solution is to let the new leader to force followers' log to duplicate its own
  - Conflicting entries in followers' logs will be ***overwritten***

# How? (I)


- Leader maintains a ***nextIndex*** for each follower
  - Index of entry it will send to that follower
- New leader sets its ***nextIndex*** to the index ***just after its last log entry***
  - 11 in the example
- Broadcasts it to all its followers



## How? (II)

- Followers that have missed some AppendEntry calls will refuse all further AppendEntry calls – *consistency check*
- Leader will decrement its nextIndex for that follower and redo the previous AppendEntry call
  - Process will be repeated until a point where the logs of the leader and the follower **match**
- Will then send to the follower all the log entries it missed

# How? (III)



- By successive trials and errors, leader finds out that the first log entry that follower (b) will accept is log entry 5
- It then forwards to (b) log entries 5 to 10

# Interesting question

- How will the leader know which log entries it can commit
  - Cannot always gather a majority since some of the replies were sent to the old leader
- Fortunately for us, any follower accepting an `AcceptEntry` RPC implicitly acknowledges it has processed all previous `AcceptEntry` RPCs

**Followers' logs cannot skip entries**

# A last observation

- Handling log inconsistencies does not require a special sub algorithm
  - Rolling back `AppendEntry` calls is enough


# Safety

- Two main issues
  - What if the log of a new leader did not contain all previously committed entries?
    - Must impose conditions on new leaders
  - How to commit entries from a previous term?
    - Must tune the commit mechanism

# Election restriction (I)

- The log of any new leader ***must*** contain all previously committed entries
  - Candidates include in their ***RequestVote*** RPCs information about the state of their log
    - *Details in the paper*
  - Before voting for a candidate, servers check that the log of the candidate is at least as up to date as their own log.
    - Majority rule does the rest

# Election restriction (II)



Two majorities of the same cluster ***must*** intersect

# Committing entries from a previous term

- A leader cannot immediately conclude that an entry from a previous term is committed even if it is stored on a majority of servers.
  - *See paper*
- Leader should never commit log entries from previous terms by counting replicas
- Should only do it for entries from the current term
- Once it has been able to do that for one entry, all prior entries are committed indirectly




# Cluster membership changes

- Not possible to do an atomic switch
  - Changing the membership of all servers at once
- Will use a two-phase approach:
  - Switch first to a transitional ***joint consensus*** configuration
  - Once the joint consensus has been committed, transition to the new configuration

# The joint consensus configuration

- Log entries are transmitted to all servers, old and new
- Any server can act as leader
- Agreements for entry commitment and elections requires majorities from both old and new configurations
- Cluster configurations are stored and replicated in special log entries

# The joint consensus configuration



# Implementations

- Two thousand lines of C++ code, not including tests, comments, or blank lines.
- About 25 independent third-party open source implementations in various stages of development
- Some commercial implementations

- A good description of how Raft works:

<http://thesecretlivesofdata.com/raft/>

# Primary Backup Replication

- Primary/backup: ensure a single order of ops:
  - Primary orders operations
  - Backups execute operations in order

# Case study: Hypervisor

*Bressoud and Schneider. Hypervisor-Based Fault Tolerance. SOSP 1995*

- Goal: fault tolerant computing
  - Banks, NASA etc. need it
  - CPUs are most likely to fail due to complexity
- Hypervisor: primary/backup replication
  - If primary fails, backup takes over
  - Caveat: assuming failure detection is perfect

# Hypervisor replicates at VM-level

- Why replicating at VM-level?
  - Hardware fault-tolerant machines were big in 80s
  - Software solution is more economical
  - Replicating at O/S level is messy (many interfaces)
  - Replicating at app level requires programmer efforts
  - Replicating at VM level has a cleaner interface (and no need to change O/S or app)
- Primary and backup execute the same sequence of machine instructions



# A Strawman design



- Two identical machines
- Same initial memory/disk contents
- Start execute on both machines
- Will they perform the same computation?


# Strawman flaws

- To see the same effect, operations must be deterministic
- What are deterministic ops?
  - ADD, MUL etc.
  - Read time-of-day register, cycle counter, privilege level?
  - Read memory?
  - Read disk?
  - Interrupt timing?
  - External input devices (network, keyboard)

# Hypervisor's architecture



Strawman replicates disks at both machines  
**Problem:** disks might not behave identically (e.g. fail at different sectors)



Hypervisor connects devices to to both machines

- Only primary reads/writes to devices
- Primary sends read values to backup
- Only primary handles interrupts from h/w
- Primary sends interrupts to backup

# Hypervisor executes in epochs

- **Challenge:** must execute interrupts at the same point in instruction streams on both nodes
- Strawman: execute one instruction at a time
  - Backup waits from primary to send interrupt at end of each instruction
  - Very slow....
- Hypervisor executes in epochs
  - CPU h/w interrupts every N instructions (so both nodes stop at the same point)
  - Primary delays all interrupts till end of an epoch
  - Primary sends all interrupts to backup

# Hypervisor failover

- If primary fails, backup must handle I/O
- Suppose primary fails at epoch  $E+1$ 
  - In Epoch  $E$ , backup times out waiting for  $[\text{end}, E+1]$
  - Backup delivers all buffered interrupts at the end of  $E$
  - Backup starts epoch  $E+1$
  - Backup becomes primary at epoch  $E+2$

# Hypervisor failover

- Backup does not know if primary executed I/O epoch  $E+1$ ?
  - Relies on O/S to re-try the I/O
- Device needs to support repeated ops
  - OK for disk writes/reads
  - OK for network (TCP will figure it out)
  - How about keyboard, printer, ATM cash machine?

# Hypervisor implementation

- Hypervisor needs to trap every non-deterministic instruction
  - Time-of-day register
  - HP TLB replacement
  - HP branch-and-link instruction
  - Memory-mapped I/O loads/stores
- Performance penalty is reasonable
  - A factor of two slow down
  - How about its performance on modern hardware?

# Caveats in Hypervisor

- Hypervisor assumes failure detection is perfect
- What if the network between primary/backup fails?
  - Primary is still running
  - Backup becomes a new primary
  - Two primaries at the same time!
- Can timeouts detect failures correctly?
  - Pings from backup to primary are lost
  - Pings from backup to primary are delayed